

# Lecture 16

## Harr Cascades

---

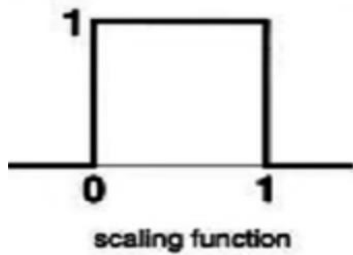
ECE 1390/2390

# Haar wavelet

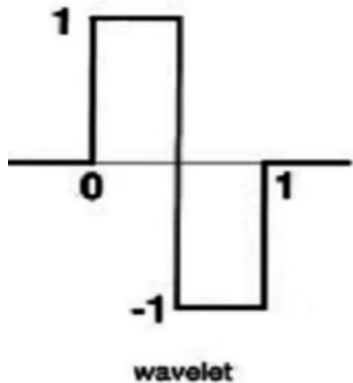
$$H_{2 \times 2} = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$$

Integrator: Low-pass filter

Differentiator: High-pass filter



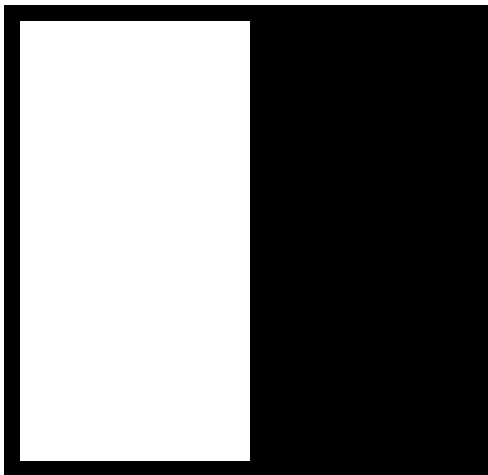
$$\phi(x) = \begin{cases} 1 & 0 \leq x < 1 \\ 0 & \text{else} \end{cases}$$



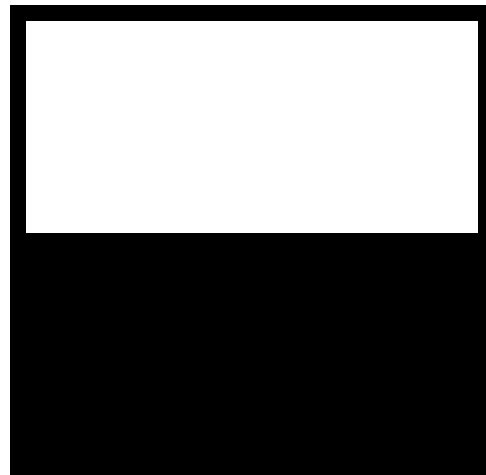
$$\psi(x) = \begin{cases} 1 & 0 \leq x < 1/2 \\ -1 & 1/2 \leq x < 1 \\ 0 & \text{else} \end{cases}$$

- Viola and Jones, "[Rapid object detection using a boosted cascade of simple features](#)", [Computer Vision and Pattern Recognition](#), 2001
  - Haar-like feature extraction kernel (not actual Haar wavelets)
  - Series of convolution kernels to extract features
  - More features than pixels
    - E.g. 24 x 24 probe has 18,000 possible rectangle arrangements
  - ` Uses only small subset of the possibilities

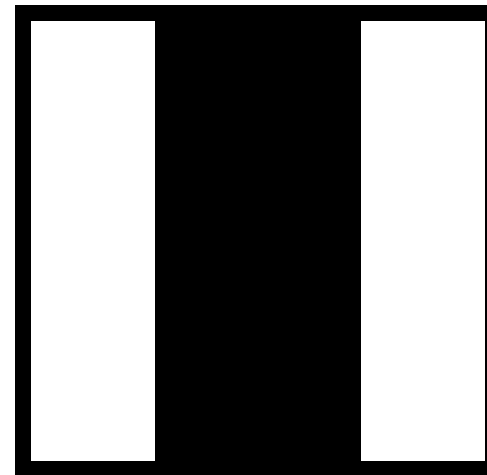
Horizontal Boundary



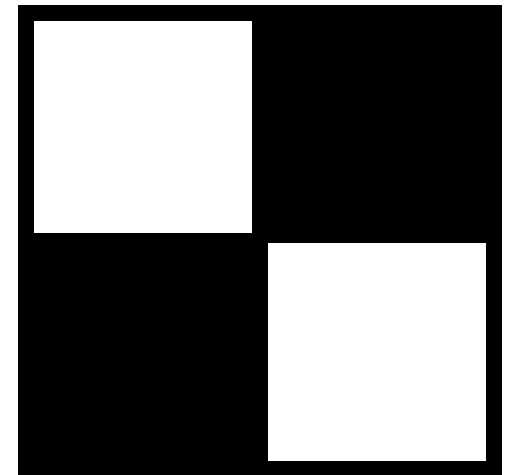
Vertical Boundary

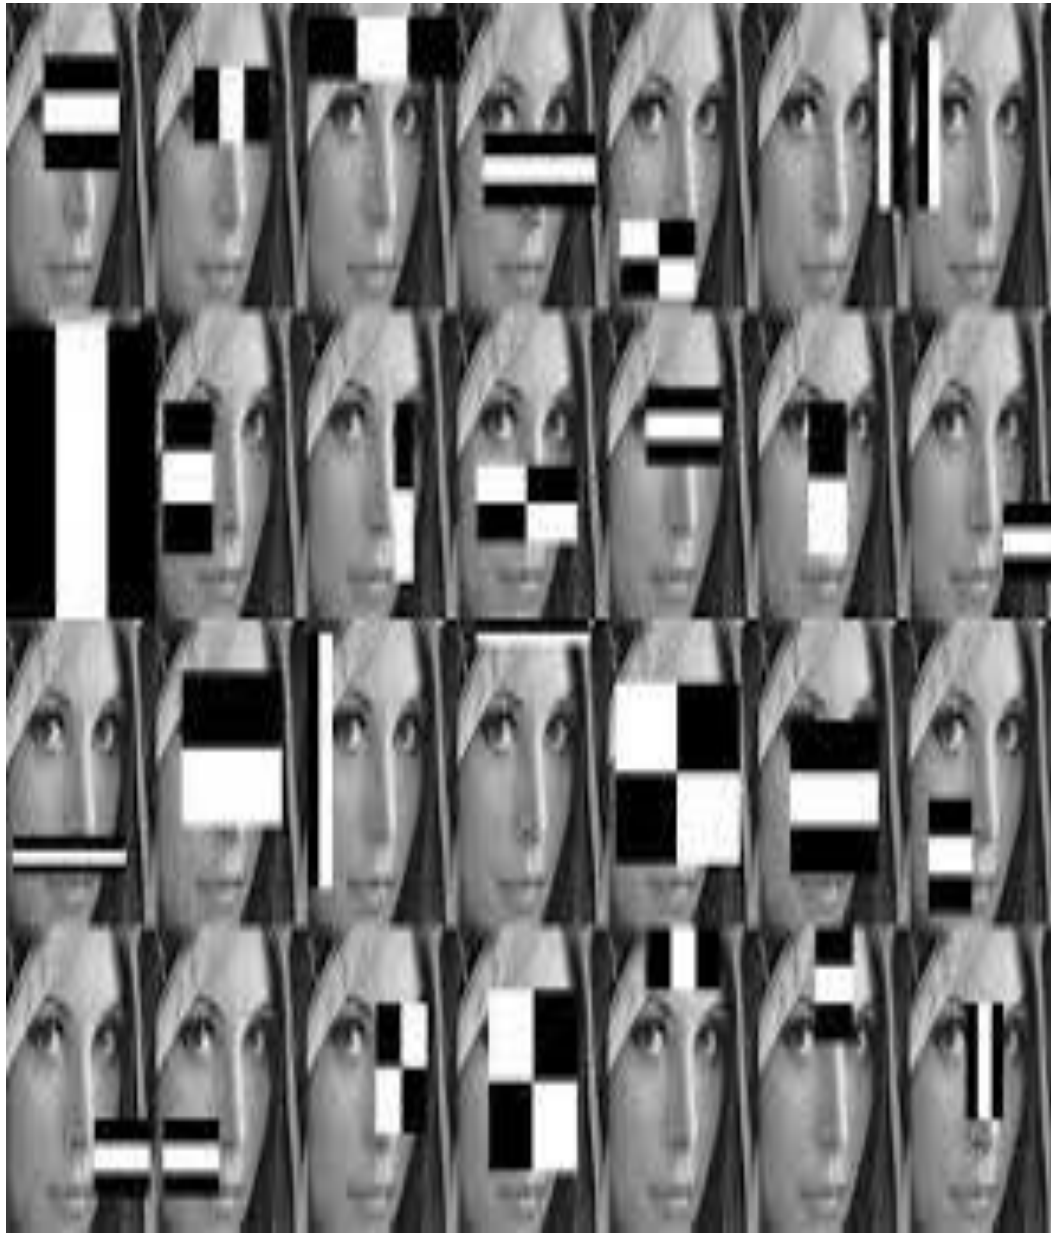


Vertical Line

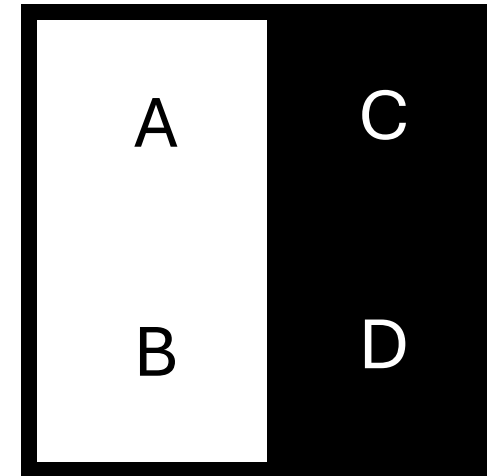


Crossing

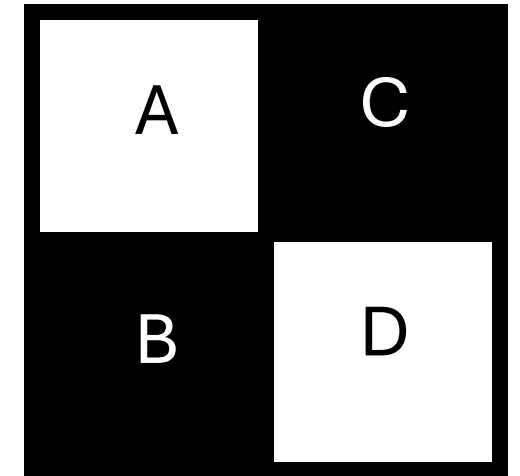




Feature Extraction:  
Sum over rectangular windows



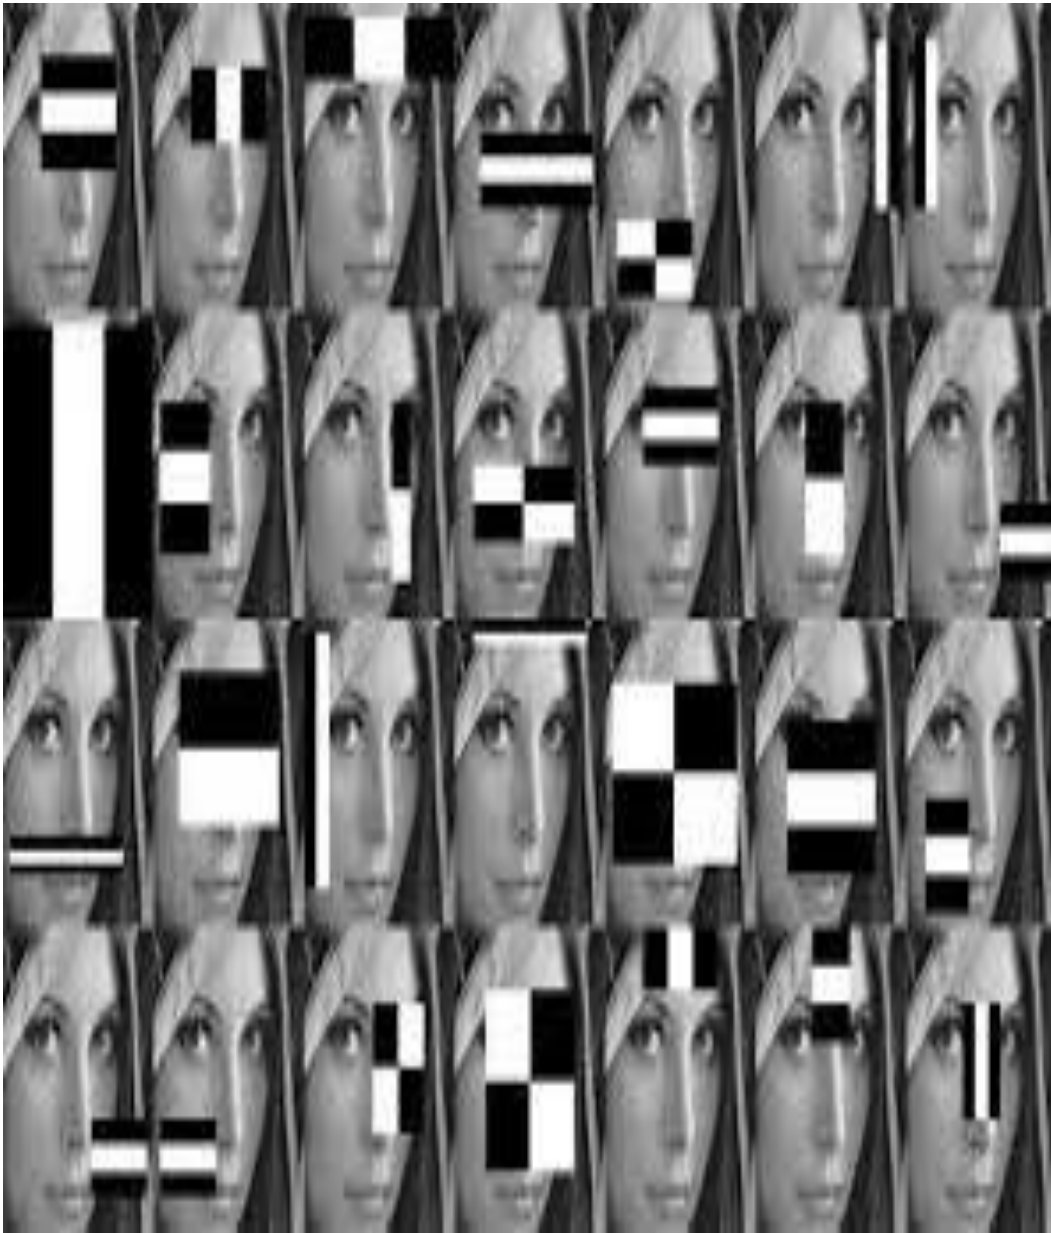
$$A+B - (C+D)$$



$$A+D - (B+C)$$

→ 28 features

Each feature is very  
weakly predictive



## Adaboost (Adaptive boosting)

$$F = \sum_i f_i(x)$$

Boasted classifier

Weak binary (-1,1) classifiers

Yoav Freund and Robert Schapire (1995), [\*A desicion-theoretic generalization of on-line learning and an application to boosting\*](#), Lecture Notes in Computer Science, Berlin, Heidelberg: Springer Berlin Heidelberg, pp. 23–37



Positive Training  
Data



Training Data



$$F = \sum_i \alpha_i \cdot f_i(x)$$

Boasted classifier

Weak binary (-1,1) classifiers

$$\begin{bmatrix} 1 \\ -1 \\ \vdots \end{bmatrix} = \begin{bmatrix} f_0(P_0) & f_1(P_0) & f_2(P_0) & \dots \\ f_0(P_1) & f_1(P_1) & f_2(P_2) & \vdots \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} \alpha_0 \\ \alpha_1 \\ \alpha_2 \\ \vdots \end{bmatrix}$$



Negative Training  
Data





Positive Training  
Data



Training Data



$$\begin{bmatrix} 1 \\ -1 \\ \vdots \end{bmatrix} = \begin{bmatrix} f_0(P_0) & f_1(P_0) & f_2(P_0) & \dots \\ f_0(P_1) & f_1(P_1) & f_2(P_1) & \vdots \\ \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} e^{\alpha_0} \\ e^{\alpha_1} \\ e^{\alpha_2} \\ \vdots \end{bmatrix}$$

Ensure that coefficients are non-negative

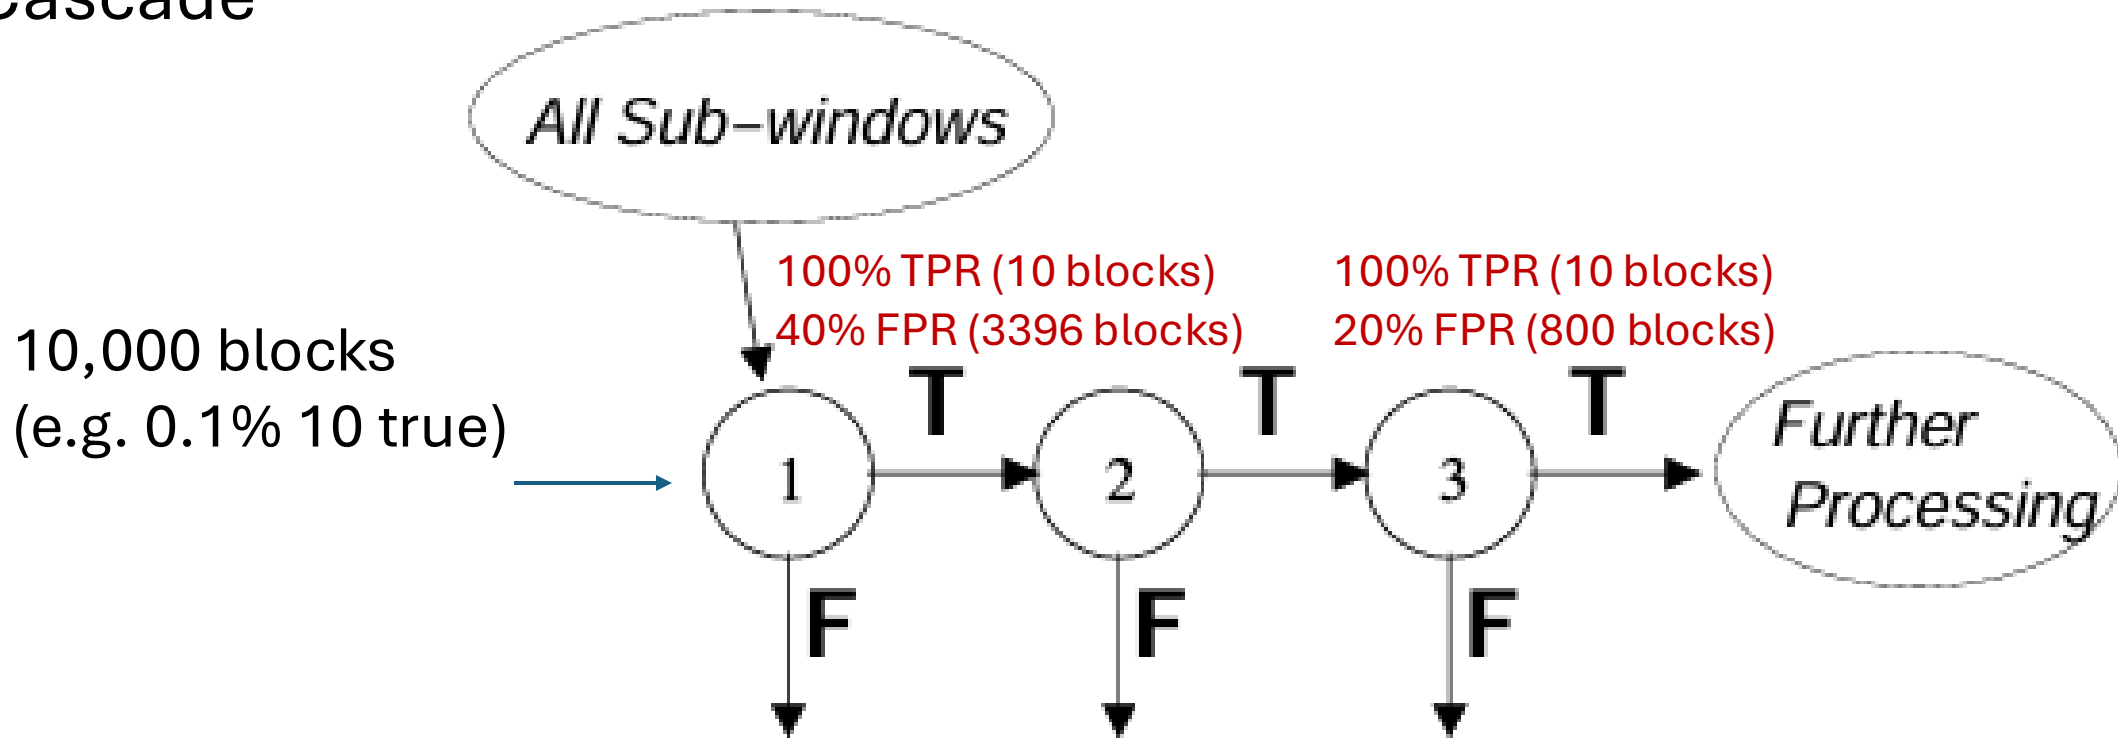
$$F = \sum_i e^{\alpha_i} \cdot f_i(x)$$

Boasted  
classifier

Weak binary  
(-1,1)  
classifiers

# Classifier Cascade

Most negatives are discarded in first stages, reducing computation



In original paper:

- 6000+ features with 38 stages
- 1, 10, 25, 25 and 50 features in the first five stages.



# Training Haar models

Good working example:

<https://github.com/bilardi/how-to-train-cascade>

Not included in python install (need to install from source code)

<https://github.com/opencv/opencv/apps/annotation>

... /createsamples

... /traincascade

# Training Haar models

Good working example:

<https://github.com/bilardi/how-to-train-cascade>

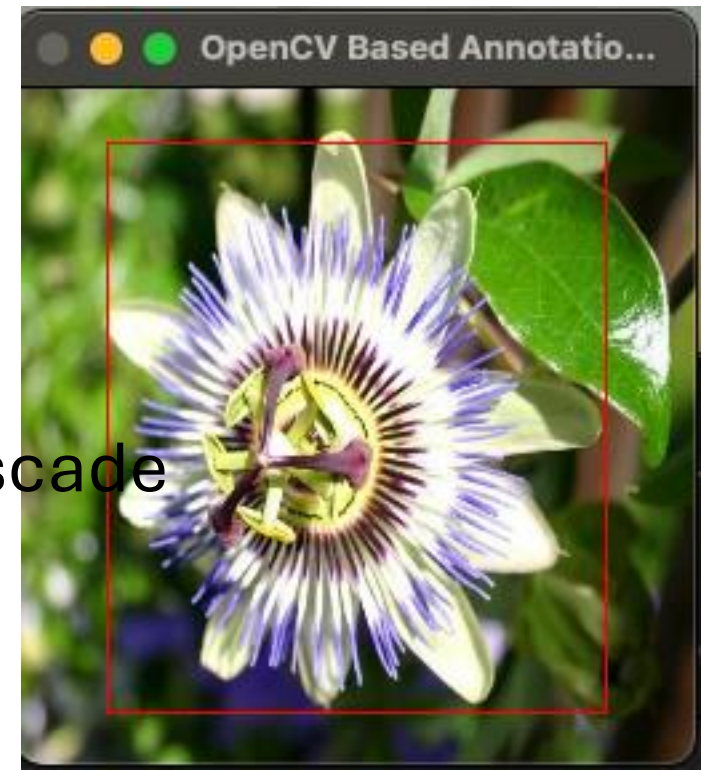
## **opencv\_annotation**

For interactive labeling of images using left mouse click followed by keystrokes

e.g. `./opencv/build/bin/opencv_annotation \`  
    `--images=/Users/theodorehuppert/Desktop/Haa/training/positives/ \`  
    `--annotations=/Users/theodorehuppert/Desktop/Haa/training/annotations.txt`

Note- if you don't give it the full pathnames, it doesn't do anything (but also gives no help info)  
Must hit "C" (red box turns green) to add annotation. Each image can have multiple annotations

- Pressing c : confirm the annotation, turning the annotation green and confirming it is stored
- Pressing d : delete the last annotation from the list of annotations (easy for removing wrong annotations)
- Pressing n : continue to the next image
- Pressing ESC : this will exit the annotation software



# Training Haar models

Good working example:

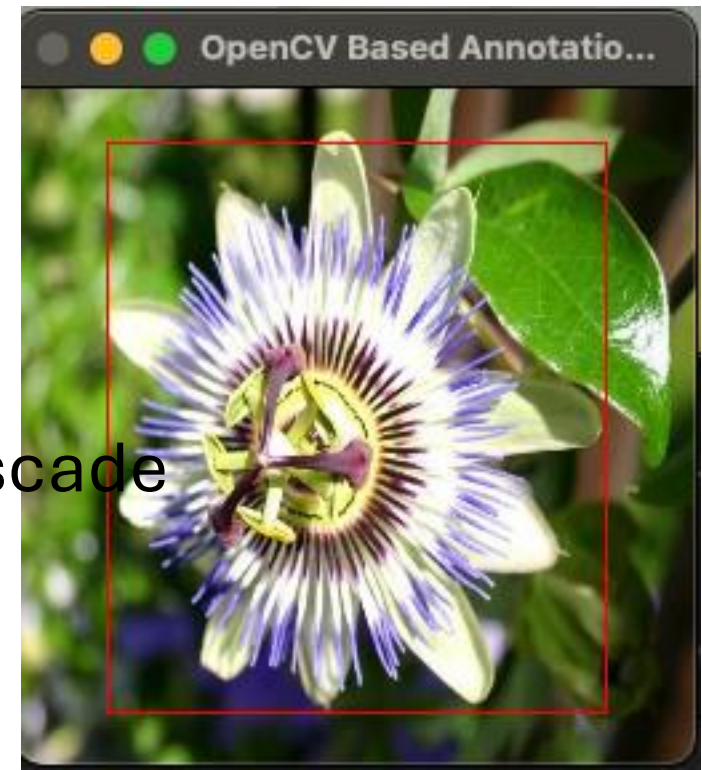
<https://github.com/bilardi/how-to-train-cascade>

## **opencv\_annotation**

For interactive labeling of images using left mouse click followed by keystrokes

e.g. `./opencv/build/bin/opencv_annotation \`  
    `--images=/Users/theodorehuppert/Desktop/Haa/training/positives/ \`  
    `--annotations=/Users/theodorehuppert/Desktop/Haa/training/annotations.txt`

```
/Users/theodorehuppert/Desktop/Haa/training/positives/image_00001.jpg 1 28 12 198 230
/Users/theodorehuppert/Desktop/Haa/training/positives/image_00002.jpg 1 45 17 194 231
/Users/theodorehuppert/Desktop/Haa/training/positives/image_00003.jpg 1 22 68 221 179
/Users/theodorehuppert/Desktop/Haa/training/positives/image_00004.jpg 1 63 31 154 206
/Users/theodorehuppert/Desktop/Haa/training/positives/image_00005.jpg 1 76 49 162 173
```



# Training Haar models

Good working example:

<https://github.com/bilardi/how-to-train-cascade>

## **opencv\_annotation**

Usage: ./opencv/build/bin/opencv\_traincascade

- data: Folder (must exist) to store output data
- vec: Samples vector created by createsamples function
- bg: Negative samples list (TXT)
- numPos: Number of samples used in training for every stage
- numNeg: “

### **Useful options**

- precalCalValBufSize: Preallocated memory (in Mb)
- precalCalIdxBufSize: Preallocated memory (in Mb)
- numThreads
- acceptanceRatioBreakValue: When to stop training model

## **Cascade Parameters**

- numStages: number of stages (default=20)
- featureType: HAAR or LBP (local binary pattern)
- w: Width of samples (must match createsamples)
- h: Height of samples (must match createsamples)
- bt: AdaBoost type. {DAB, RAB, LB, GAB}
- minHitRate: Desired positive hit rate for each stage
- maxFalseAlarmRate: Max false positives for each stage
- mode: BASIC- only vertical HAAR. ALL—also 45degree rotated

# Training Haar models

Good working example:

<https://github.com/bilardi/how-to-train-cascade>

## **opencv\_traincascade**

Usage: ./opencv/build/bin/opencv\_traincascade

- data: Folder (must exist) to store output data
- vec: Samples vector created by createsamples function
- bg: Negative samples list (TXT)
- numPos: Number of samples used in training for every stage
- numNeg: “

### **Useful options**

- precalCalValBufSize: Preallocated memory (in Mb)
- precalCalIdxBufSize: Preallocated memory (in Mb)
- numThreads
- acceptanceRatioBreakValue: When to stop training model

## **Cascade Parameters**

- numStages: number of stages (default=20)
- featureType: HAAR or LBP (local binary pattern)
- w: Width of samples (must match createsamples)
- h: Height of samples (must match createsamples)
- bt: AdaBoost type. {DAB, RAB, LB, GAB}
- minHitRate: Desired positive hit rate for each stage
- maxFalseAlarmRate: Max false positives for each stage
- mode: BASIC- only vertical HAAR. ALL—also 45degree rotated

- `git clone https://github.com/bilardi/how-to-train-cascade`
- `cd how-to-train-cascade`
- `find ./positive_images -iname "*.jpg" > positives.txt`
- `find ./negative_images -iname "*.jpg" > negatives.txt`

# Note in the git repo, info.dat has already been created for you

- `../opencv/build/bin/opencv_createsamples -info info.dat -num 37 -w 80 -h 80 -vec samples.vec`
- `../opencv/build/bin/opencv_traincascade -data classifier \`  
    `-vec samples.vec -bg negatives.txt -numStages 50 \`  
    `-minHitRate 0.999 -maxFalseAlarmRate 0.5 -numPos 37 \`  
    `-numNeg 16 -w 80 -h 80 -precalcValBufSize 1024`  
    `-precalcIdxBufSize 1024 -featureType HAAR -mode BASIC`



# Training Haar models

Good working example:

<https://github.com/bilardi/how-to-train-cascade>

## Suggestions

- HAAR classifiers work for “rigid, textured objects”. (e.g. faces onward, but faces side-view would need a different filter)
- Want to have around 1000-2000 positive images
- Want to have about 2:1 positive:negative images
- Total false discovery is  $(\text{maxFalseAlarmRate [per stage]})^{\text{number stages}}$
- Total true positive rate is:  $(\text{minHitRate [per stage]})^{\text{number stages}}$ 
  - minHitRate: Desired positive hit rate for each stage
  - maxFalseAlarmRate: Max false positives for each stage

# Training Haar models

## **Useful list of a bunch of free (large) curated datasets:**

<https://imerit.net/blog/28-free-image-datasets-for-computer-vision-all-pbm/>

<https://image-net.org/>. [requires registration with a .edu email]. (Basis for ResNet)

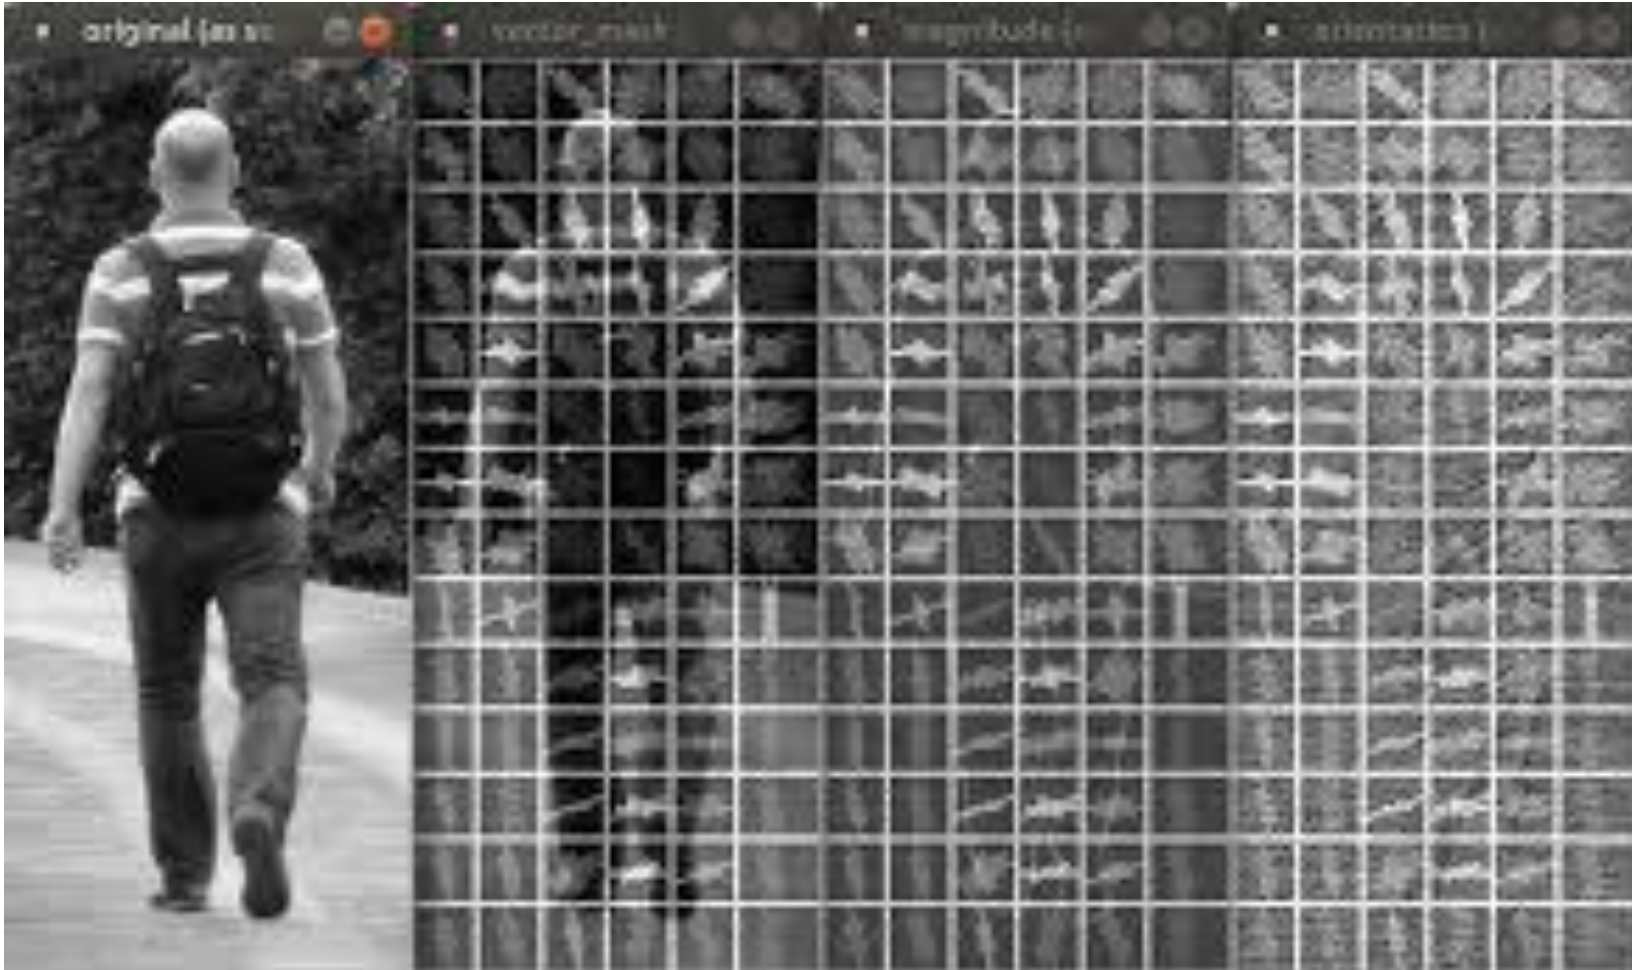
<https://www.robots.ox.ac.uk/~vgg/data/>

## **Trained models**

**<https://github.com/opencv/opencv/data/haarcascades>**

eyes, front faces, profile faces, fullbody, upper body, lower body, smile, Russian license plates (???)

# HOG (Histogram of oriented gradients)



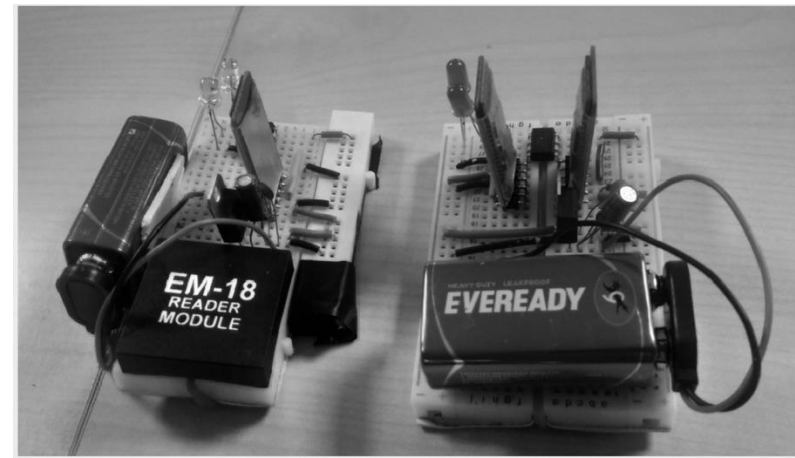
# Sobel filter

Low-pass  
filter

High-pass  
filter

$$\begin{array}{c} \left| \begin{array}{c} 1 \\ 2 \\ 1 \end{array} \right| \\ \text{1D Gaussian} \\ \text{Filter} \end{array} * \begin{array}{c} \overline{\overline{1 \ 0 \ -1}} \\ \text{x - Derivative} \end{array} = \begin{array}{c} \left| \begin{array}{ccc} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{array} \right| \\ \text{Sobel - X} \end{array}$$

$$\begin{array}{c} \left| \begin{array}{c} 1 \\ 0 \\ -1 \end{array} \right| \\ \text{y - Derivative} \end{array} * \begin{array}{c} \overline{\overline{1 \ 2 \ 1}} \\ \text{1D Gaussian} \\ \text{Filter} \end{array} = \begin{array}{c} \left| \begin{array}{ccc} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{array} \right| \\ \text{Sobel - Y} \end{array}$$

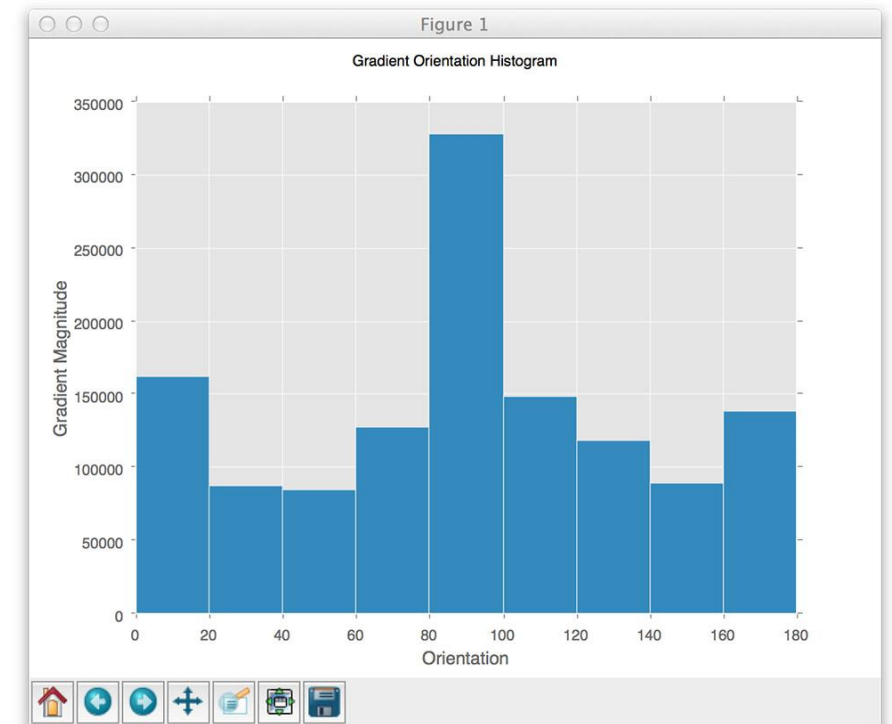


$$G = \sqrt{G_X^2 + G_Y^2}$$

$$\theta = \tan^{-1} \left( \frac{GX}{GY} \right)$$

# HOG (Histogram of oriented gradients)

Each block position provides  
9 feature values (#bins used in this case)



# HOG (Histogram of oriented gradients)

## SVM Classifier

Given positive/negative samples (same as Haar cascade),  
compute the HOG features for each sample  $\rightarrow$  N features

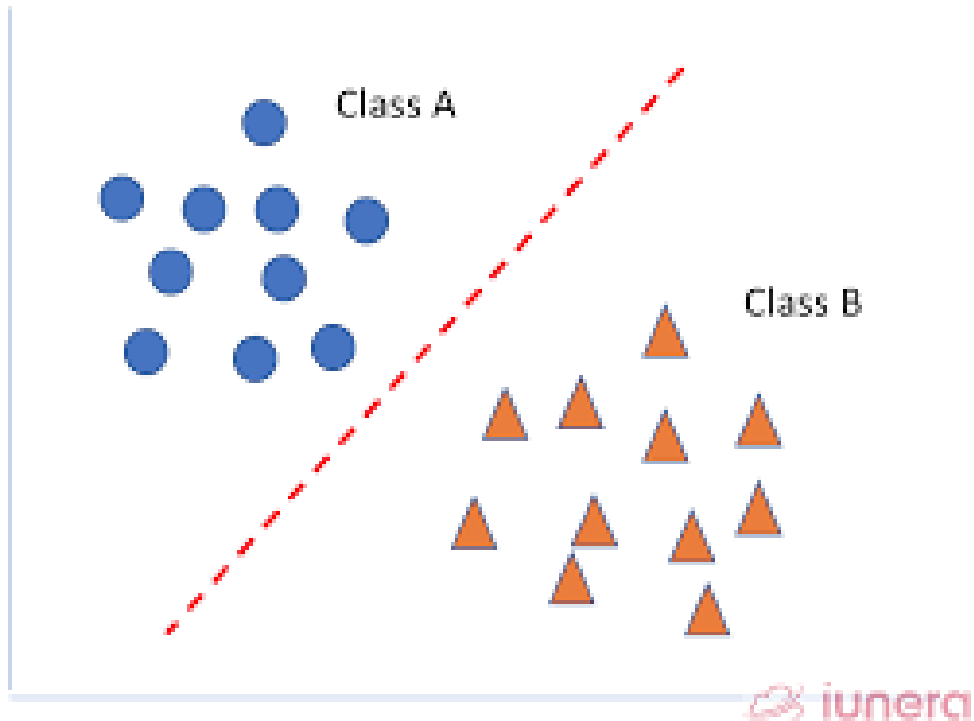
+1 : Positive

0: Negative

Construct feature space matrix (X) <# features x examples>  
and label vector ( $Y = \in [0,1]$ )

Find hyperplane (w) that minimizes

$$\arg \min_w w^T \cdot X - y = 0$$





# HOG (Histogram of oriented gradients)

```
hog = cv2.HOGDescriptor()  
hog.setSVMDetector(cv2.HOGDescriptor_getDefaultPeopleDetector())
```

# Detect people in the image

```
(rects, weights) = hog.detectMultiScale(imgray, winStride=(4, 4),  
                                         padding=(8, 8), scale=1.05)
```

